

Reese Ludwick

reesepludwick@gmail.com | (304)-886-1850 | <https://www.linkedin.com/in/reese-ludwick-024333286/> |
<https://github.com/reese8272>

SUMMARY

Applied AI engineer who ships LLM and agent applications to production end-to-end — RAG pipelines, multi-agent LangGraph systems, and custom MCP tool servers deployed for real clients. Strengths in applied AI system design, prompt engineering, and model/output evaluation, with hands-on delivery across FastAPI + Python backends, PostgreSQL/Redis, and async task pipelines. Focused on turning state-of-the-art LLM tooling into dependable, scalable software that works in front of real users.

EXPERIENCE

Software Associate

Remote

Cognizant

Dec 2025 – Present

- Build GenAI and agentic AI systems with LangChain, LangGraph, and Python across R&D and client deliverables.
- Design and ship RAG pipelines and MCP server integrations for production workflows.
- Scope, prototype, and deliver AI data-analysis tools alongside cross-functional teams.

Freelance AI Engineer & Full-Stack Developer

Remote

WorldCovers / APMC

May 2026 – Present

- Paid contract engineer on production web properties and the WorldCovers platform.
- Ship features and automation against a 14-table PostgreSQL schema using Django REST + React/TypeScript.
- AI-augmented engineering (Claude Code) under compressed timelines, with direct stakeholder coordination.

GenAI Developer Contractor

Remote

Cognizant

Jun – Dec 2025

- Built GenAI POCs with LangChain, LangGraph, and Python.
- Ran prompt engineering and model evaluation to optimize LLM outputs across business use cases.
- Built custom FastMCP servers and external MCP integrations with AWS for dynamic agent tool access.

IT Student Helper

James Rumsey Technical Institute

2019 – 2021

- Resolved 300+ support tickets; led a Wi-Fi overhaul cutting latency 25%; completed 100+ AV/network setups.

PROJECTS

CreatorClip / autoclip.studio — Python, FastAPI, Celery, pgvector, WhisperX, Voyage AI, Anthropic API, ffmpeg, YouTube OAuth

Live SaaS (Stripe billing), shipped end-to-end. Async Celery workers run WhisperX transcription, Voyage + pgvector embeddings for semantic retrieval, and ffmpeg rendering. Custom creator-“DNA” module tailors output; YouTube OAuth for publishing. Production LLM/ML orchestration at scale; ~79% test coverage.

CFO Agent (in active development) — Python, FastAPI, LangGraph, PostgreSQL, Redis, Anthropic API

FastAPI + LangGraph agent with PostgreSQL persistence and Redis caching over imported financial data (CSV/OFX import + encrypted vault). Hybrid context architecture; multi-tenant-ready design. Applied AI reasoning over financial data.

Agentic Software Development Team — LangChain, LangGraph, Amazon Bedrock (Cognizant R&D)

Multi-agent LangGraph system (codegen / testing / security / supervision) automating the build lifecycle on Bedrock and OSS ML backends.

Customer Call Data Analysis & Insights Engine — Python, NLP, Sentence Transformers (Cognizant R&D)

NLP clustering of call data into semantic issue groups; agentic AI for call consolidation and executive reporting.

TECHNICAL SKILLS

Languages: Python, SQL

AI / ML: LangChain, LangGraph, RAG, MCP Servers, Amazon Bedrock, Sentence Transformers, TensorFlow, pgvector, Anthropic API, Voyage AI, WhisperX, prompt engineering, model/output evaluation

Backend / Infra: FastAPI, Django REST, PostgreSQL, Redis, Celery, Docker

Frontend: React / TypeScript

Data: Pandas, NumPy, Jupyter

Tooling / Cloud: Linux, Ubuntu, Git, GitHub, VS Code, Claude Code, Azure Fundamentals

EDUCATION

B.S. Computer Science — Shepherd University (Aug 2021 – May 2025) — Major GPA 3.43

CERTIFICATIONS

TensorFlow Developer Specialization (DeepLearning.ai) · Linux + Git (Boot.dev) · AZ-900 · GitHub Copilot · Generative AI Context Engineering (Cognizant)